


Machine learning

Introduction

Small overview of Machine learning

Overview of machine learning history



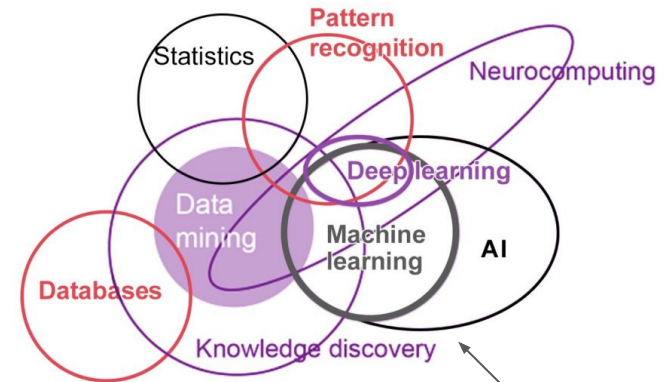
Development of computers and software: new possibilities!

Boom from the combination of: 1) computer power, 2) data and 3) ML methods

History of Science: new instruments -> new breakthroughs

From the blog:

<https://jaylatta.net/history-of-ai-from-winter-to-winter/>




Quite wrong but gives a rough idea.
What is correct?

Early days

- 1805: linear regression or method of least squares (Legendre)

Data points (x, y)

Line $y = ax + b$



From **data points**: find the line which **minimize** the sum of squared errors

The value for new points can be **predicted**.

In bold: central concepts in ML

Early days

Note: Machine learning development goes along with computer development

- 1805: linear regression or method of least squares (Legendre)
- 1943: McCulloch and Pitts: first mathematical model of neural nets
- 1950: Alan Turing, “the Turing test” (the imitation game), a test to evaluate if a machine is “intelligent”
- 1951: k-nearest neighbors
- 1957: Rosenblatt, the perceptron
- 50-60s: Bayesian methods
- 60-70s: clustering
- 1969: Minsky and Papert, end of the perceptron and beginning of the first AI winter, the XOR problem

Early days, the perceptron

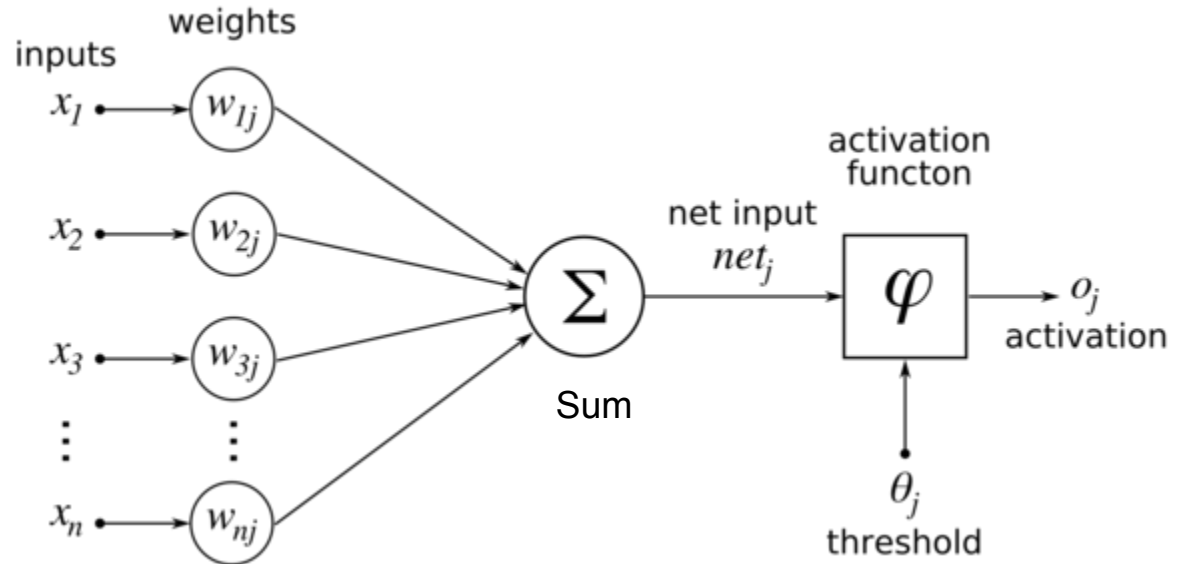






Figure 4.8 Illustration of the Mark 1 perceptron hardware. The photograph on the left shows how the inputs were obtained using a simple camera system in which an input scene, in this case a printed character, was illuminated by powerful lights, and an image focussed onto a 20×20 array of cadmium sulphide photocells, giving a primitive 400 pixel image. The perceptron also had a patch board, shown in the middle photograph, which allowed different configurations of input features to be tried. Often these were wired up at random to demonstrate the ability of the perceptron to learn without the need for precise wiring, in contrast to a modern digital computer. The photograph on the right shows one of the racks of adaptive weights. Each weight was implemented using a rotary variable resistor, also called a potentiometer, driven by an electric motor thereby allowing the value of the weight to be adjusted automatically by the learning algorithm.

The Perceptron and classification




The Perceptron create a separation line (solve a linear classification problem)


the XOR problem



A	B	Q
0	0	0
0	1	1
1	0	1
1	1	0



OR




XOR

The XO

The Perceptron create a separation line (solve a linear classification problem)



THE PERCEPTRON CAN'T
LEARN THE XOR FUNCTION



IT IS NOT LINEARLY
SEPARABLE, WE
NEED MORE RESEARCH

First AI winter 1974-1980

Results below expectations


Research fund cutted

but developments continued at a slower pace

1979: Neocognitron (for vision, inspired convolutional neural nets)

Clustering

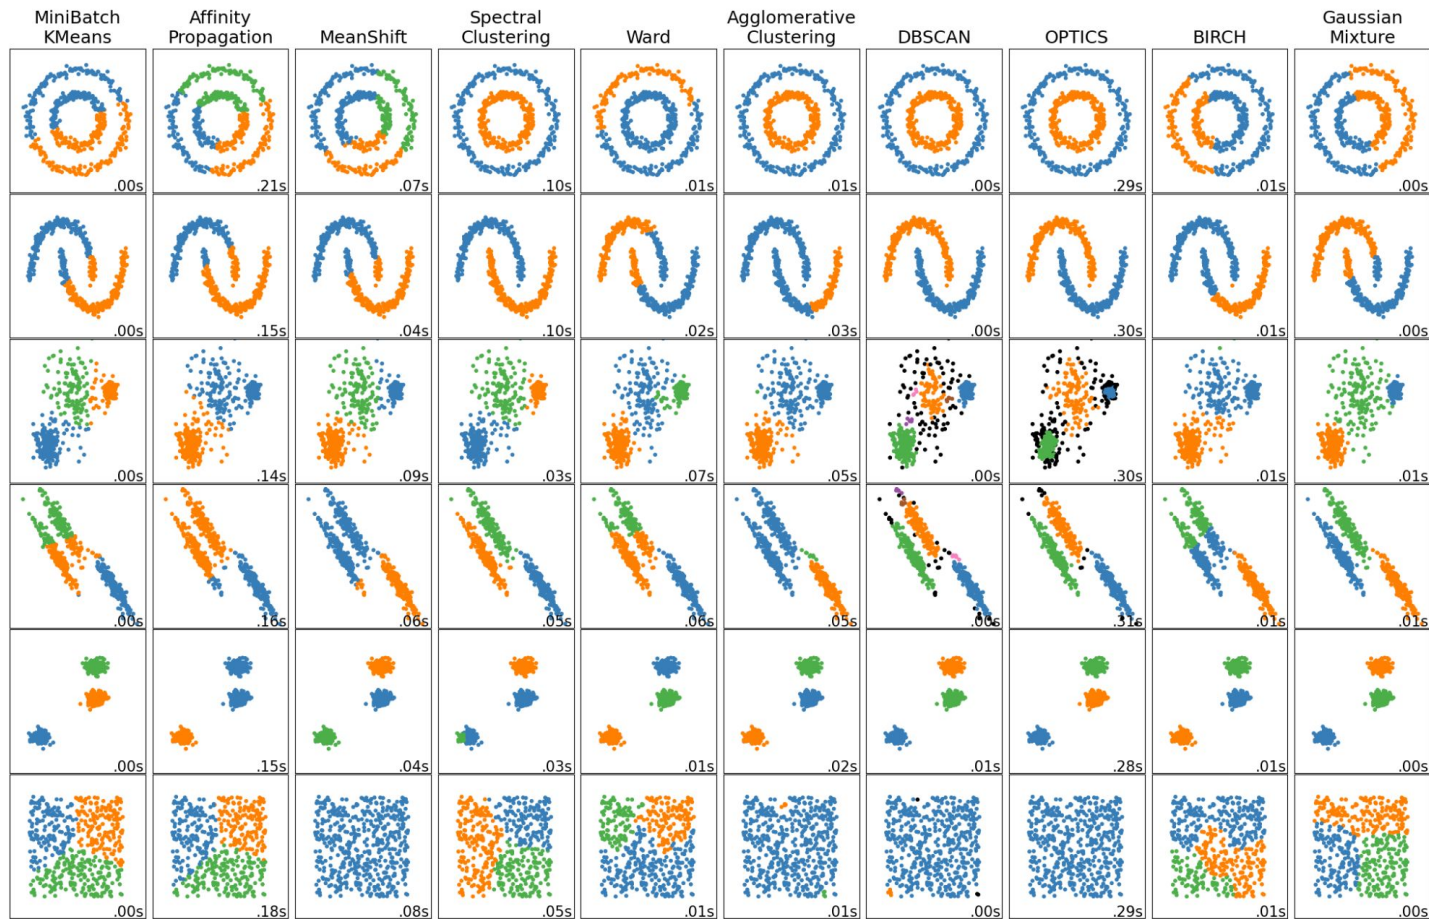
Principle, group close data points together



- most methods do not use any label (unsupervised)
- some may use a partial labelling (semi-supervised)



Clustering

From <https://scikit-learn.org/stable/modules/clustering.html>



Machine learning boom 80-87


- More powerful computers, programming languages
- Decision trees in their modern form
- Rediscovering backpropagation
- Developing several layers of neurons, multilayer perceptron



Second AI winter 1987- early 2000s

1988: universal approximation theorem. Multilayer NN can approximate any function f , $y=f(x)$. (f must satisfy a few general conditions)

1991: the vanishing gradient problem (in multilayer neural networks)



Recent developments

1995: support vector machines, random forest, gradient boosting

1997: Deep Blue beats the world champion at chess

90s: Graphical Processing Units (GPUs) appear and are progressively integrated in computers and video games

2009: publication discussing the benefits of GPUs for machine learning

2012 **the breakthrough of deep learning**: Imagenet classification and deep learning, Alexnet

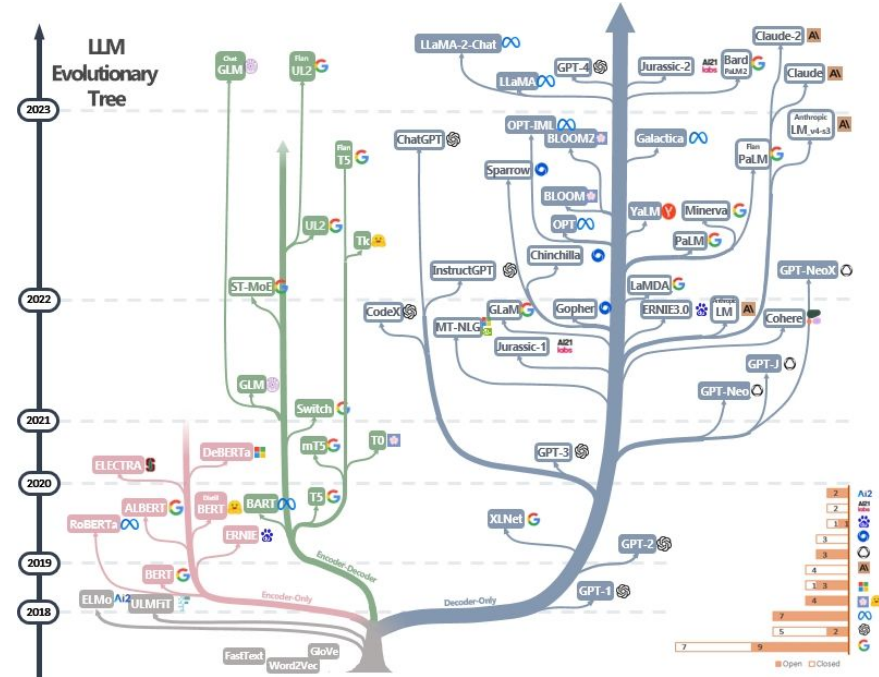
2014: GANs, generative adversarial networks.

2016: AlphaGo beats best player at Go

2017: Transformer architecture (basis for BERT, chatGPT + many others)

2021: DALL-E, imagen, stable diffusion: AI-generated images

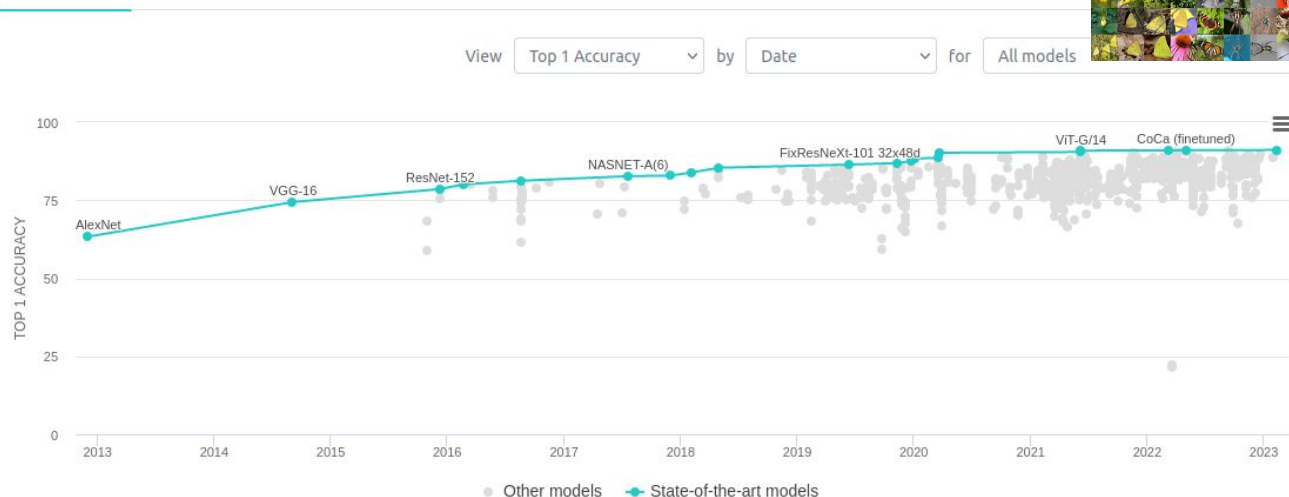

2022: chatGPT



2012 breakthrough on imagenet

- A large dataset Imagenet
- First use of GPUs for ML
- New set of machine learning concepts and tricks

Deep learning





from <https://paperswithcode.com/sota/image-classification-on-imagenet>

Why AI goes so fast now?

1) GPUs

Graphical Processing Units (graphic card)

- Designed to compute many small operations in parallel
- graphic cards in computers
- are now used for deep learning
- very efficient for matrix multiplication



Meta will have 350,000 of Nvidia's fastest AI GPUs by end of year, buying AMD's MI300, too

News

By Anton Shilov published January 19, 2024

Meta expects to have 600,000 H100 GPU-equivalent AI horsepower by end of 2024.

2) Data

Machine learning needs data

Imagenet (2009) <https://www.image-net.org/>

- 1000 object classes
- 1 million images

Laion-5B (2022) <https://laion.ai/>

- 5 billion image-text pairs
- scrapped from the web

The Pile <https://pile.eleuther.ai/>

An 800GB Dataset of Diverse Text for Language Modeling

what is the size of all the Harry Potter books?

<https://www.kaggle.com/datasets/moxxis/harry-potter-lstm>



3) Open-source

- GAFAM develop free and open source tools for deep learning (Pytorch, Tensorflow...)
- Active open-source community
 - free access to software
 - free access to course and learning
 - easy to contribute

Non-AI people are using AI tools: generating images, generating text with chatGPT

An active open-source community

Active open-source community

- easy to contribute with Git and Github
- scikit-learn
- ML tutorials
- Kaggle
- Huggingface provides coded, pre-trained models, ready to use

<https://scikit-learn.org/>



<https://github.com/>



<https://www.kaggle.com/>



<https://huggingface.co/>



AI going too fast?

AI and images, generative models

New AI methods with many applications to images:
image generation from text prompt, denoising, inpainting,
super-resolution...



<https://openai.com/blog/dall-e-introducing-outpainting>



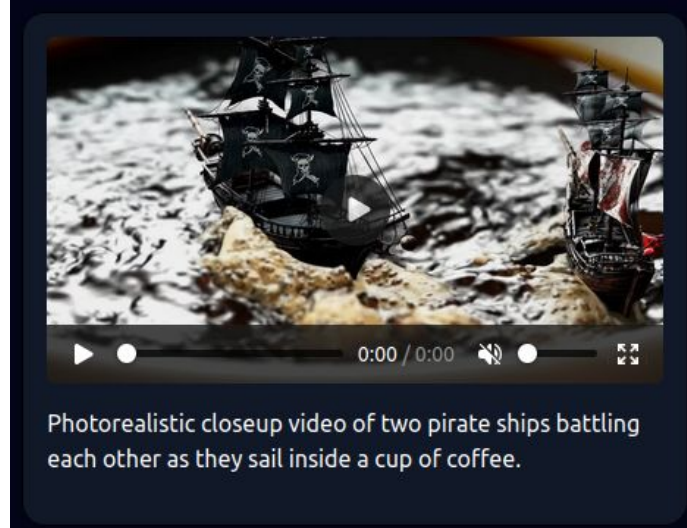
generative AI



Images
Twitter search: #Midjourney #Pope

Videos

<https://sora.aitubo.ai/>



Photorealistic closeup video of two pirate ships battling each other as they sail inside a cup of coffee.

<https://x.com/EccentrismArt/status/1823059492520788342>

AI generated music

<https://suno.com/>

Generating clichés, gender and diversity bias

<https://openai.com/blog/reducing-bias-and-improving-safety-in-dall-e-2>

A portrait of a software engineer

Generate



Before mitigation

After mitigation

Awful AI

<https://github.com/daviddao/awful-ai>


daviddao / **awful-ai** Public

Sponsor Watch 296


<> Code Issues 12 Pull requests 7 Discussions Actions Projects Security Insights

master 2 branches 1 tag

Go to file Add file <> Code


 **daviddao** Added 'forensic sketching' in awful AI use case ...

ec3dca7 on Feb 6 67 commits

 .github

Create FUNDING.yml

4 years ago

 README.md

Added 'forensic sketching' in awful AI use case

2 months ago

☰ README.md

Awful AI


Awful AI is a curated list to track *current* scary usages of AI - hoping to raise awareness to its misuses in society

Artificial intelligence in its current state is [unfair](#), [easily susceptible to attacks](#) and [notoriously difficult to control](#). Often, AI systems and predictions [amplify existing systematic biases](#) even when the data is balanced. Nevertheless, more and more concerning uses of AI technology are appearing in the wild. This list aims to track *all of them*. We hope that *Awful AI* can be a platform to spur discussion for the development of possible preventive technology (to fight back!).


You can [cite the list](#) and raise more awareness through Zenodo.


DOI [10.5281/zenodo.5855972](https://doi.org/10.5281/zenodo.5855972)


About


 Awful AI is a curated list to track current scary usages of AI - hoping to raise awareness

twitter.com/dwddao

 Readme


 6.6k stars

 296 watching

 227 forks


[Report repository](#)

Releases 1

 **Awful AI - 2021 Edition** Latest

on Jan 16, 2022

Sponsor this project

 **daviddao** David Dao

More info on AI (BBC article)

FUTURE NOW | ARTIFICIAL INTELLIGENCE

The A-Z of AI: 30 terms you need to understand artificial intelligence

<https://www.bbc.com/future/article/20230717-what-you-should-know-about-artificial-intelligence-from-a-z>

More history

https://en.wikipedia.org/wiki/Timeline_of_machine_learning

Timeline for AI and images: <https://www.fabianmosele.com/ai-timeline>

<https://www.historyofdatascience.com/>